

Conceptual Engineering Using Large Language Models

Bradley P. Allen

2023-12-16

How are philosophy and AI alike?

For the past 400 years, the domain of philosophy has been shrinking. ... Physics, geology, chemistry, economics, biology, anthropology, sociology, meteorology, psychology, linguistics, computer science, cognitive science. Each of those subject matters was a part of philosophy a mere 400 years ago.

Kevin Scharp. Philosophy as the study of defective concepts. In Conceptual engineering and conceptual ethics, pages 396–416. Oxford University Press, 2020.

... the promoters of AI may not have delivered exactly what they announced ... but what they delivered is astounding. Many breakthroughs in computer science ... can directly be traced to work in AI. Part of the problem is a phenomenon that I heard John McCarthy himself describe: *“As soon as it works, no one calls it AI anymore.”*

Bertrand Meyer. "John McCarthy". <https://cacm.acm.org/blogs/blog-cacm/138907-john-mccarthy/fulltext>, 2011. Accessed: 2023-03-03.

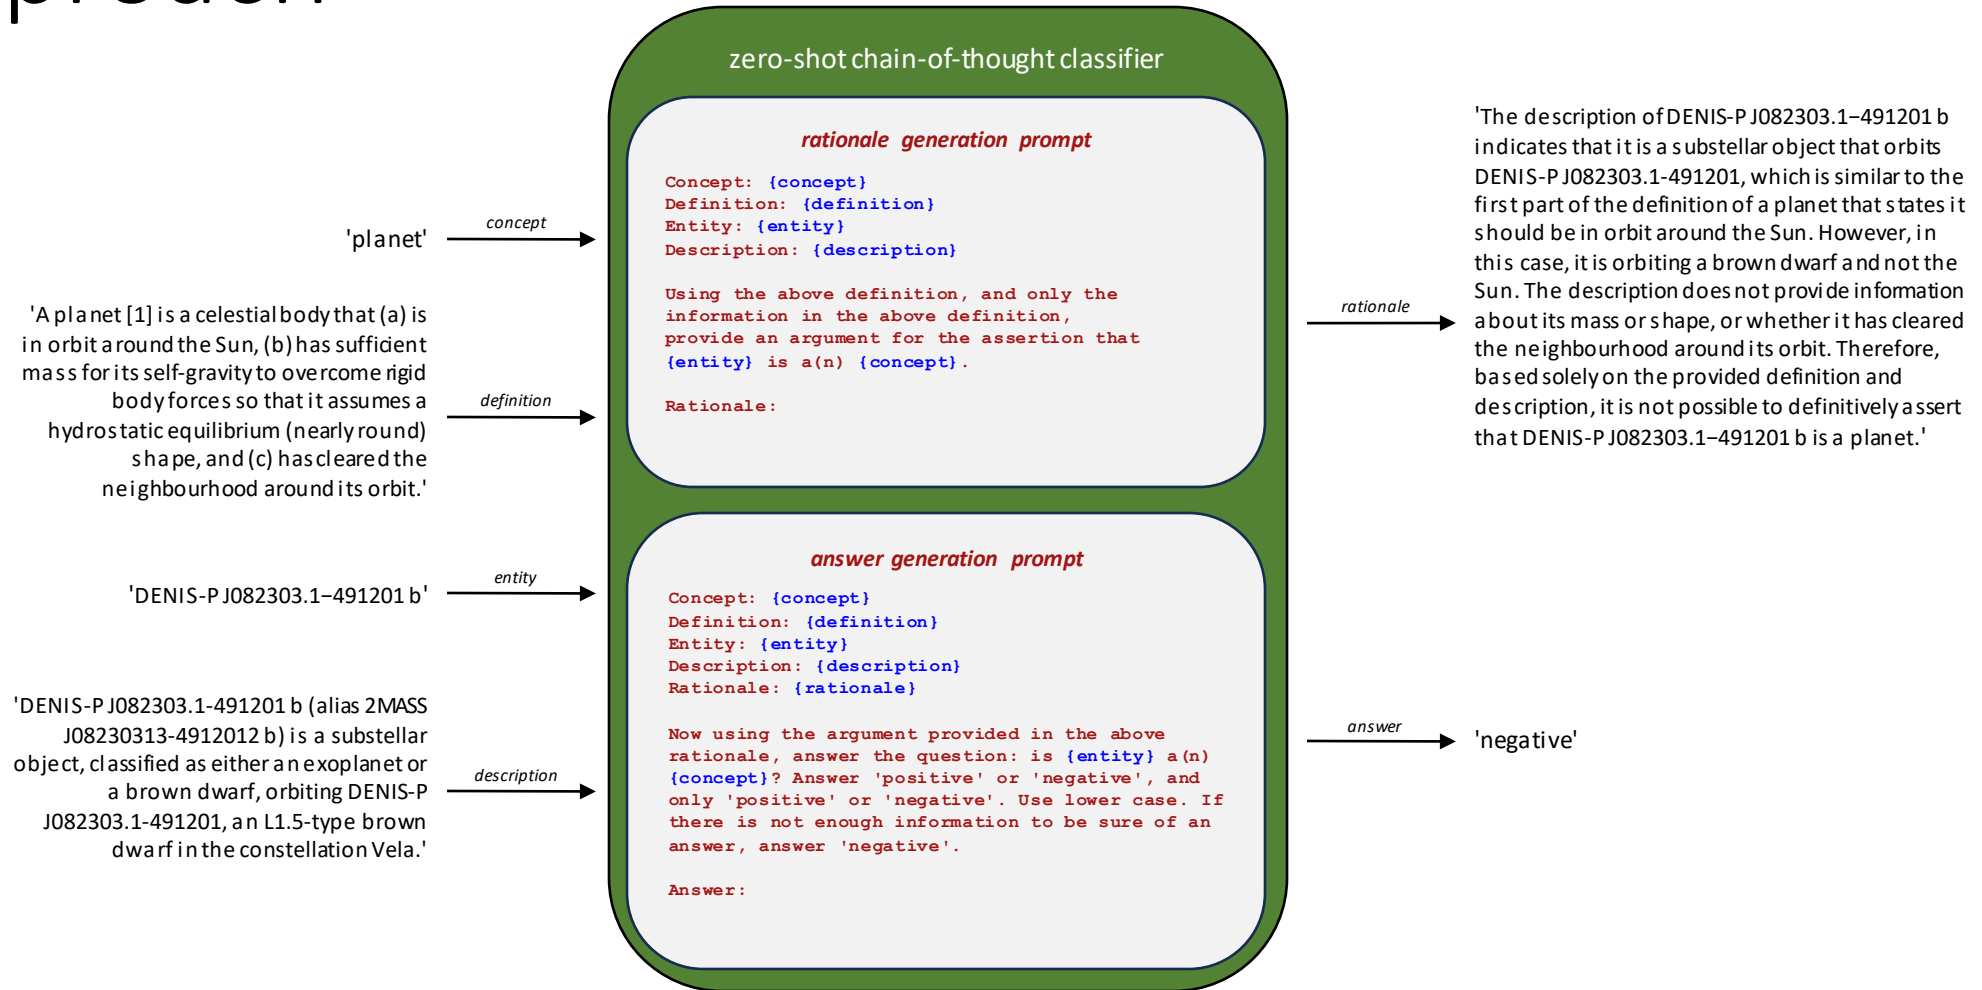
How are conceptual engineering and knowledge engineering alike?

- Philosophers have different ideas about the best way to do conceptual engineering (CE), but it usually involves expressing and analyzing concepts using natural language
- Knowledge engineering (KE) focuses on transforming knowledge expressed in natural language into a structured formal language suitable for automated reasoning
- We hypothesize that CE and KE are instances of the same activity: *refining and transforming knowledge to achieve a normative goal*
- We believe that large language models (LLMs) create an opportunity for a unification of these two practices

Background

- An important question for a theory of conceptual engineering (CE) is the nature of its targets, i.e., "what conceptual engineers are (or should be) trying to engineer"
- Nado proposed as targets *classification procedures* (CPs), defined as abstract 'recipes' which sort entities "into an 'in'-group and an 'out'-group"
- We build on this idea by defining a method that uses prompt engineering of LLMs to implement CPs
- We then evaluate CPs built using this method by leveraging a knowledge graph (KG) as a source of positive and negative examples of elements in the extension of a concept

Approach



Evaluation

- Obtain natural language definitions for a concept
- Sample positive and negative examples of the concept from Wikidata
- For each example, retrieve a summary of its Wikipedia page to use as its description
- Apply the classification procedure for a given definition of the concept to each example, yielding a classification of the example
- Compute a confusion matrix from the classifications
- Review false positives/negatives to determine if a given error arises from the concept's definition or from the entity's description

Datasets

PLANET

Definitions:

- OED
- Haslanger
- Homosaurus

Positive examples: 50 instances with sex or gender (P21) either female (Q6581072) or trans woman (Q1052281)

Negative examples: 50 instances with sex or gender either male (Q6581097), non-binary (Q48270), or trans man (Q2449503)

WOMAN

Definitions:


- OED
- IAU 2006-08-16
- IAU 2006-08-24

Positive examples: 50 instances (P31) of planet (Q634)


Negative examples: 50 instances of substellar object (Q3132741) that are not instances of planet

Performance metrics

PLANET (N = 100)



WOMAN (N = 100)



concept	definition	Cohen's kappa	F1 macro	FN	FP
PLANET	IAU 2006-08-24	0.96	0.98	1	1
	IAU 2006-08-16	0.94	0.97	1	2
	OED	0.92	0.96	1	3
WOMAN	Homosaurus	0.96	0.98	0	2
	Haslanger	0.94	0.97	2	1
	OED	0.92	0.96	2	2

Findings

- Rationales generated by the classification procedures were sound
- Answers were faithful to their rationales
- Rationales frequently contained statements about issues with definitions or descriptions
- For PLANET, most errors were false positives relating to trans-Neptunian objects
- For WOMAN:
 - Homosaurus performed best, possibly because it is the most inclusive definition
 - Haslanger had two false negatives, the rationales of which ascribed to entity descriptions lacking evidence of systematic subordination

Discussion

- We've shown how a CE project can incorporate an empirical, data-driven activity using LLMs and KGs
- It raises the possibility that generative AI assistants could support philosophers in the conduct of CE projects
- Using CPs to align natural language concepts with knowledge graphs can be of practical value for data management and governance, which contributes new success conditions for CE
- However, the use of LLMs as they exist today raises transparency, reproducibility and safety concerns
- Further work is needed to evaluate our method with respect to these issues, with a specific focus on evaluating explanation faithfulness

Conclusion

- We constructed a conceptual engineering target as a computational artifact
- We then applied it to provide an empirical evaluation of the target
- We view this as an initial step in an investigation of the potential utility of large language models in the practice of conceptual engineering



UNIVERSITEIT
VAN AMSTERDAM

INDE lab

Thank you!

GitHub repository: <https://github.com/bradleypallen/zero-shot-classifiers-for-conceptual-engineering>